

C++ to C Transpiler: Architecture and Motivation

Version: 1.0

Date: 2025-12-26

Purpose: Understand the 3-stage pipeline architecture and why it was designed this way

Table of Contents

1. Overview
 2. The 3-Stage Pipeline
 3. Why This Architecture?
 4. Key Design Principles
 5. Alternative Approaches Considered
 6. Benefits in Practice
-

Overview

The C++ to C transpiler converts modern C++ code into clean, readable, formally-verifiable C99 code. The architecture is built around a **3-stage pipeline** that separates concerns and optimizes for **output code quality** over implementation simplicity.

The Core Insight

There are three common approaches to transpilation, and one novel approach (ours):

1. **Direct text emission from source AST** - Walk C++ AST, emit text strings
2. **LLVM IR/bytocode approach** - Lower to LLVM IR, emit from bytecode
3. **LLM-based transpilation** - Use GPT/Claude to translate source code
4. **Intermediate AST approach** (our novel approach) - C++ AST → **separate C AST** → text

For production-quality, human-readable, verifiable C code, approaches 1-3 fail:

Approach 1 (Direct text emission):

- ✗ Messy output with incorrect precedence
- ✗ Missing parentheses causing bugs
- ✗ Thousands of edge cases to handle manually
- ✗ Unreadable generated code

Approach 2 (LLVM IR):

- ✗ Loses high-level structure (names, types, control flow)
- ✗ Output is machine-like, not human-readable
- ✗ Cannot preserve source semantics for verification
- ✓ Works for JIT/VM targets (Emscripten, WASM)
- ✓ Good for performance-critical backends

Approach 3 (LLM-based transpilation):

- ✗ **Hallucinations** - Invents APIs, functions, syntax that don't exist
- ✗ **Inconsistency** - Same input produces different outputs
- ✗ **Context limits** - Cannot process large codebases coherently
- ✗ **No guarantees** - Cannot prove correctness or completeness
- ✗ **Missing edge cases** - Fails on corner cases not in training data
- ✗ **Non-deterministic** - Cannot reproduce builds
- ! **Poor context engineering** - Struggles with cross-file dependencies
- ✓ Useful for prototypes, learning, simple one-off conversions
- ✓ Good for human-in-the-loop assisted migration

Our novel solution (Approach 4): Build a **separate, complete C AST** (using Clang's C node types) and let Clang's battle-tested printer handle code generation.


What makes this novel:


- Most transpilers emit text directly (Approach 1) or use LLVM IR (Approach 2)
- Clang's TreeTransform API modifies AST in-place, doesn't create separate output AST
- We build an entirely new AST using **only C-compatible node types** (RecordDecl, FunctionDecl, CallExpr, etc.)
- Then we leverage Clang's printer which has solved all the hard problems over 15+ years
- No known C++ to C transpiler uses this architecture (as of 2025)

Why this matters for our goals:

- Preserves high-level semantics (unlike LLVM IR)
 - Clean, readable output (unlike direct emission)
 - Suitable for Frama-C verification (unlike LLM approach)
 - Deterministic and reproducible (unlike LLM approach)
 - Debuggable with source mapping
 - Handles edge cases systematically (unlike LLM approach)
-

The 3-Stage Pipeline





Stage 1: C++ AST Generation

What: Parse C++ source code using Clang frontend

Input: C++ source files

Output: Fully-resolved C++ AST with type information, template instantiations, semantic analysis

Why: Leverage Clang's world-class C++ parser instead of building our own

Stage 2: C++ AST → C AST Translation

What: Walk C++ AST and build equivalent C AST using handler chain

Input: C++ AST from Stage 1

Output: Pure C AST with only C-compatible nodes

Why: Make ALL translation decisions here, create proper C structures

Stage 3: C Code Emission

What: Print C AST to text using Clang's DeclPrinter/StmtPrinter

Input: C AST from Stage 2

Output: Clean, formatted C source code

Why: Leverage 15+ years of production-tested code generation

Why This Architecture?

Problem: Direct Text Emission is Fragile

Naive Approach (What We DON'T Do):

```
// BAD: Emit text directly from C++ AST
void emitExpression(Expr* E) {
    if (BinaryOperator* BO = dyn_cast<BinaryOperator>(E)) {
        emitExpression(BO->getLHS());
        OS << " " << BO->getOpcodeStr() << " ";
        emitExpression(BO->getRHS());
        // PROBLEM: What about precedence?
        // PROBLEM: What about parentheses?
        // PROBLEM: What about edge cases?
    }
}
```

Issues:

- ✗ Incorrect operator precedence: `a + b * c` becomes `(a + b) * c`
- ✗ Missing parentheses cause subtle bugs
- ✗ Edge cases pile up (pointer-to-member, overloaded operators, etc.)
- ✗ Thousands of lines of brittle string concatenation
- ✗ Hard to verify output correctness

Solution: Intermediate C AST

Our Approach (What We DO):

```
// GOOD: Build C AST nodes
Expr* translateExpression(Expr* E) {
    if (BinaryOperator* BO = dyn_cast<BinaryOperator>(E)) {
        Expr* lhs = translateExpression(BO->getLHS());
        Expr* rhs = translateExpression(BO->getRHS());
        return builder.createBinaryOp(BO->getOpcode(), lhs, rhs);
        // Clang's printer handles precedence automatically!
    }
}
```

Benefits:

- Clang's printer handles precedence, parentheses, formatting
- 15+ years of edge cases already fixed
- Zero maintenance burden for code generation
- Output is clean and verifiable

The Numbers: Output Quality

Metric	Direct Emission	Intermediate AST	Improvement
Lines per exception handler	46 lines	11 lines	4.2x cleaner
Code size	100%	20%	80% reduction
Frama-C verification time	100%	10-20%	5-10x faster
Implementation LOC	1,400-2,300	2,000-3,200	+40% code

Trade-off (Pre-AI Era, 2015-2023): We accepted 40% more implementation code for 80% cleaner output.

Trade-off (AI Era, 2024+): With AI-assisted development (Claude Code, Cursor, Copilot), implementation effort is now comparable to direct emission while retaining 80% cleaner output.

Rationale: Generated code is read/debugged/verified far more often than implementation code. With AI handling boilerplate, we get superior output quality without developer productivity cost.

Key Design Principles

1. Separation of Concerns

Each stage has ONE clear responsibility:

Stage	Responsibility	Does NOT Do
Stage 1	Parse C++	✗ Translation decisions
Stage 2	Decide C representation	✗ Text generation

Stage	Responsibility	Does NOT Do
Stage 3	Format text output	✗ Translation logic

Golden Rule: Stage 2 decides WHAT to emit, Stage 3 decides HOW to format it.

2. Testability

Each stage can be unit tested independently:

```
// Test Stage 2 WITHOUT running Stage 3
TEST(TranslationTest, ClassToStruct) {
    // Given: C++ class in AST
    CXXRecordDecl* cppClass = ...;

    // When: Translate to C
    RecordDecl* cStruct = translator.translate(cppClass);

    // Then: Verify C AST structure (no text output needed!)
    EXPECT_EQ(cStruct->getName(), "MyClass");
    EXPECT_EQ(cStruct->fields().size(), 2);
}
```

3. Battle-Tested Components

Don't Reinvent the Wheel:

- **Clang's Parser** → Used by Apple, Google, Microsoft
- **Clang's Printer** → 15+ years in production
- **PNaCl Exception Handling** → Used by Chrome Native Client
- **Itanium ABI RTTI** → Standard used by GCC, Clang, ICC

Result: Stand on the shoulders of giants.

4. Frama-C First

Optimize for formal verification:

Runtime Library Approach:

```
// Instead of inlining 46 lines of exception handling:
try_block() {
    // 46 lines of complex setjmp/longjmp code
}

// We call runtime library:
try_block() {
    cxx_frame_push(&frame); // 11 lines total
    if (setjmp(frame.jmpbuf) == 0) {
        // try body
    } else {
        // catch handler
    }
    cxx_frame_pop(&frame);
}
```

Benefits:

- Verify runtime library ONCE with Frama-C
- Generated code just calls verified functions
- 5-10x easier to verify

Alternative Approaches Considered

Comparison of Transpilation Approaches

Here's how the three common approaches plus our novel approach compare for our use case:

Aspect	Direct Text Emission	LLVM IR Approach	LLM-Based	Separate C AST (Novel)
Input	C++ AST	LLVM Bytecode	Source Text	C++ AST
Output Quality	✗ Poor (precedence bugs)	✗ Machine-like	⚠ Varies (non-deterministic)	✓ Clean, readable
Preserves Structure	⚠ Partial	✗ Lost	⚠ Sometimes	✓ Yes

Aspect	Direct Text Emission	LLVM IR Approach	LLM-Based	Separate C AST (Novel)
Variable Names	✓ Preserved	✗ Lost to SSA	✓ Usually preserved	✓ Preserved
Control Flow	✓ Preserved	✗ Basic blocks	⚠ May change	✓ Preserved
Correctness Guarantees	⚠ Manual testing	✓ Semantic preserving	✗ None (hallucinations)	✓ AST-level correctness
Deterministic Output	✓ Yes	✓ Yes	✗ No	✓ Yes
Handles Edge Cases	✗ Manual effort	✓ Systematic	✗ Training-dependent	✓ Systematic
Frama-C Verification	⚠ Difficult	✗ Incompatible	✗ Unreliable	✓ Optimized
Implementation Effort (2025)	⚠ Medium	⚠ Medium	✓ Low (API calls)	✓ Medium (AI-assisted)
Implementation Effort (Pre-AI)	⚠ Medium	⚠ Medium	N/A	✗ High (verbose)
Maintenance Burden	✗ High	✓ Low	⚠ Model-dependent	✓ Low
Best For	Quick prototypes	JIT/VM/WASM	Learning, assisted migration	Production transpilers

Note on Implementation Effort (2025):

The intermediate AST approach was previously considered "high effort" due to verbose Clang API calls. **With AI-assisted development in 2025**, this is no longer true:

- ✓ **AI generates boilerplate:** CNodeBuilder helper functions written by AI (Claude Code, Cursor, Copilot)
- ✓ **AI handles Clang API complexity:** 50-line node creation → AI-generated in seconds
- ✓ **AI assists with visitor patterns:** RecursiveASTVisitor implementation scaffolded automatically
- ✓ **Focus on high-level logic:** Developers specify WHAT to translate, AI handles HOW

- **This project itself:** Built with AI assistance - validates the approach!

Result: Implementation effort is now **medium** (comparable to other approaches) while retaining all quality benefits.

Example: Simple C++ Function

```
// Input C++
class Calculator {
    int value;
public:
    int add(int x) { return value + x; }
};
```

Approach 1 - Direct Text Emission (Buggy):

```
// Output - Missing parentheses, wrong precedence
int Calculator_add(struct Calculator *this, int x) {
    return this->value + x; // OK for this case
    // But complex expressions become: a + b * c + d
    // Should be: (a + b) * (c + d) ← Bug!
}
```

Approach 2 - LLVM IR (Machine-Like):

```
; Output LLVM IR
%class.Calculator = type { i32 }

define i32 @Calculator_add(%class.Calculator* %this, i32 %x) {
entry:
    %value.addr = getelementptr inbounds %class.Calculator, %class.Calculator* %this, i32 0
    %0 = load i32, i32* %value.addr
    %1 = add nsw i32 %0, %x
    ret i32 %1
}
; Converting back to C loses readability
```

Approach 3 - Intermediate AST (Ours - Clean):

```
// Output C - Clean, correct, readable
struct Calculator {
    int value;
};

int Calculator_add(struct Calculator *this, int x) {
    return this->value + x; // Clang printer guarantees correctness
}
```

Alternative 1: TreeTransform API

What: Use Clang's TreeTransform API to modify C++ AST in-place

Why We Rejected It:

- Clang documentation: "TreeTransform does not support adding new nodes well"
- Requires 50+ lines of boilerplate per node type
- Cannot inject statements or create new variables easily
- Still need a printer afterward
- **Score:** 4.1/10 (see docs/architecture/prototype-comparison.md)

Decision: Direct C AST generation scores 9.2/10

Alternative 2: Direct Text Emission

What: Walk C++ AST and emit text strings directly

Why We Rejected It:

- Fragile operator precedence handling
- Thousands of edge cases to handle manually
- Unreadable output with missing parentheses
- Impossible to verify correctness
- **Result:** 3-5x worse output quality

Decision: Intermediate AST + Clang printer

Alternative 3: LLVM IR Lowering

What: Use LLVM backend to lower C++ to LLVM IR, then emit C from IR

Valid Use Cases for LLVM IR:

- Emscripten (C/C++ → WebAssembly) - Machine code target
- JIT compilation - Runtime optimization
- Cross-platform VMs - Portable bytecode
- Optimization-heavy backends - Leverage LLVM passes

Why We Rejected It for Human-Readable C:

- Loses high-level semantics (class names, variable names, structure)
- Output is SSA form (Single Static Assignment), not readable C
- Control flow becomes basic blocks and phi nodes
- Cannot preserve source structure for code review
- Not suitable for formal verification (Frama-C needs readable C)
- **Example:** `int add(int a, int b) { return a + b; }` becomes:

```
define i32 @add(i32 %a, i32 %b) {
entry:
    %0 = add nsw i32 %a, %b
    ret i32 %0
}
```

Converting back to C loses names, structure, and readability

Decision: Stay at AST level, preserve high-level structure for human consumption

Note: LLVM IR is an excellent choice for machine-targeted transpilation (WebAssembly, JVM bytecode, native code), just not for our goal of human-readable, verifiable C.

Alternative 4: LLM-Based Transpilation

What: Use Large Language Models (GPT-4, Claude, etc.) to translate source code via prompts

Valid Use Cases for LLM Transpilation:

- Quick prototypes and proof-of-concepts
- Learning and educational purposes
- Human-in-the-loop assisted migration (developer reviews each output)
- One-off conversions of small codebases
- Migration planning and strategy development

Why We Rejected It for Production Transpiler:

- **✗ Hallucinations:** Invents non-existent functions, APIs, or syntax
 - Example: `std::vector::fast_sort()` (doesn't exist)
 - Creates plausible-looking but incorrect code
- **✗ Non-deterministic:** Same input produces different outputs
 - Run 1: Uses approach A
 - Run 2: Uses approach B (equally valid but different)
 - Cannot reproduce builds
- **✗ Context window limits:** Cannot coherently process large files
 - 100K token limit = ~30K lines of code
 - Loses consistency across large codebases
 - Forgets earlier decisions
- **✗ Poor context engineering:** Struggles with cross-file dependencies
 - Missing type definitions from headers
 - Incorrect function signatures
 - Lost namespace context
- **✗ No correctness guarantees:** Cannot prove output matches input semantics
 - Might compile but behave differently
 - Subtle bugs in edge cases
 - Cannot verify with formal methods
- **✗ Training data bias:** Fails on patterns not in training set
 - Modern C++20 features may be under-represented
 - Domain-specific patterns (embedded, safety-critical) may be mishandled
 - Non-standard compiler extensions likely wrong
- **✗ Cost and latency:** API calls per file add up
 - Large project = thousands of API calls
 - Latency makes iteration slow
 - Cost scales with codebase size

Example of LLM Hallucination:

```
// Input C++
std::vector<int> vec = {1, 2, 3};
vec.shrink_to_fit();
```

LLM might output (WRONG):

```
// Hallucinated – realloc doesn't exist in this context!
int* vec = malloc(3 * sizeof(int));
vec = realloc_shrink(vec, 3); // Function doesn't exist!
```

Correct AST-based output:

```
// Correct – systematic translation
struct vector_int {
    int* data;
    size_t size;
    size_t capacity;
};
vector_int_shrink_to_fit(&vec); // Calls runtime function
```

Decision: AST-based approach provides deterministic, verifiable, production-quality output.

Note: LLMs are powerful tools for assisted development and code exploration, but not suitable for automated production transpilation where correctness and reproducibility are critical.

Benefits in Practice

1. Clean, Readable Output

C++ Input:

```
class Point {
    int x, y;
public:
    Point(int x, int y) : x(x), y(y) {}
    int getX() const { return x; }
};
```

Generated C (Actual Output):

```

struct Point {
    int x;
    int y;
};

void Point_ctor(struct Point *this, int x, int y) {
    this->x = x;
    this->y = y;
}

int Point_getX(const struct Point *this) {
    return this->x;
}

```

Note: No extra parentheses, correct formatting, human-readable.

2. Debugging with #line Directives

Generated code maps back to original C++ source:

```

#line 42 "example.cpp"
void MyClass_method(struct MyClass *this) {
#line 43 "example.cpp"
    int x = this->value;
#line 44 "example.cpp"
    return process(x);
}

```

Compiler errors reference **original C++ line numbers**, not generated C.

3. Zero Maintenance for Code Generation

Clang's printer is maintained by LLVM community:

- Automatic bug fixes
- New C standard features
- Performance improvements
- Edge case handling

We never touch the printer code.

4. Formal Verification

Frama-C can verify:

```
/*@ requires this != NULL;
 @ ensures this->x == x && this->y == y;
 */
void Point_ctor(struct Point *this, int x, int y);
```

Runtime library verified ONCE, reused everywhere.

The Bottom Line

Why 3 Stages?

1. **Stage 1 (Clang Parse):** Let Clang handle the hard part (C++ parsing)
2. **Stage 2 (Translation):** We control C representation (our expertise)
3. **Stage 3 (Clang Print):** Let Clang handle the hard part (text generation)

Why Our Novel Intermediate C AST Approach?

What's different from other transpilers:

- Most C++ transpilers emit text directly (fragile, buggy)
- Some use LLVM IR (unreadable, loses structure)
- We build a **separate, complete C AST** - to our knowledge, first of its kind

Why this works:

Pre-AI Era (2015-2023): We traded implementation complexity for output quality:

- +40% more implementation code (verbose Clang API)
- -80% generated code size
- 3-5x cleaner output
- 5-10x easier verification

AI Era (2024+): We get the best of both worlds:

- AI handles verbose boilerplate → implementation effort comparable to direct emission

- Still get 80% cleaner output
- Still get 5-10x easier verification
- Zero maintenance burden (Clang printer)

For a transpiler, output quality is everything. With AI assistance, we no longer sacrifice developer productivity to achieve it.

Why This Matters

This architecture enables:

- Clean C code suitable for code review
- Formal verification with Frama-C
- Debugging with original source locations
- Safety-critical embedded systems use
- Long-term maintainability

Result: A production-quality transpiler, not a research prototype.

Why We Believe This is Novel

Research conducted:

- Analyzed historical transpilers (Cfront, Comeau C++)
- Studied modern tools (Emscripten, emmtrix eCPP2C)
- Examined Clang tooling (TreeTransform, clang-tidy, clang-refactor)
- Reviewed academic literature on source-to-source translation

Findings:

- **Cfront (1983-1993):** Likely emitted text directly, no public AST-building documentation
- **Emscripten:** Uses LLVM IR → asm.js/WASM path, not C AST
- **emmtrix eCPP2C:** Commercial tool, no public architecture details, likely proprietary approach
- **Clang TreeTransform:** Modifies AST in-place, doesn't create separate C AST
- **Academic tools:** Focus on analysis/refactoring, not complete transpilation

Our contribution:

- First documented approach to build **separate C AST from C++ AST**
- Systematically use Clang's C node types (not C++ nodes) for output
- Leverage printer as zero-maintenance backend

- Optimize for formal verification (Frama-C)

Caveat: Other proprietary transpilers may use similar techniques but haven't published their architectures. If you know of prior art, please let us know!

Meta: This Project Validates AI-Assisted Development

This transpiler was built with AI assistance (**Claude Code**), validating our architectural choice:

Pre-AI Concern (2015-2023):

- "Intermediate AST requires too much boilerplate"
- "Clang API is too verbose for manual coding"
- "Direct text emission is faster to implement"

AI-Era Reality (2024+):

- **CNodeBuilder helpers:** AI generated 500+ lines in minutes
- **Visitor patterns:** AI scaffolded RecursiveASTVisitor implementations
- **AST node creation:** AI handles 50-line Clang API calls automatically
- **Documentation:** AI helped write this comprehensive architecture doc
- **Tests:** AI assists with test generation for edge cases


Result: We implemented the "high-effort" approach in the time it would have taken to do "quick" direct emission, but got production-quality output.


Lesson: In 2025, **architectural decisions should prioritize output quality over implementation simplicity**, because AI handles implementation complexity.

Further Reading

- **Implementation Details:** [docs/architecture/01-pipeline-architecture.md](#)
- **Design Decision Rationale:** [docs/architecture/architecture-decision.md](#)
- **Prototype Comparison:** [docs/architecture/prototype-comparison.md](#)
- **Complete Architecture:** [docs/ARCHITECTURE.md](#)

Visual Summary





Document Status: Living Document

Maintained By: Project Team

Questions? See [Contributing Guide](#)